Zooplankton Hierarchical Classification

CNN-RNN Sequence of Tokens Approach

Philip Loewen - University of Toronto, Department of Statistical Sciences

The Problem

What are we classifying?

12 classes of images with 8 species of zooplankton from microscopy images spanning cladocera, rotifers, copepods.

Why is it difficult?

Classes share a natural taxonomic hierarchy which classifiers ignore. This structure improves model interpretability and why errors occur.

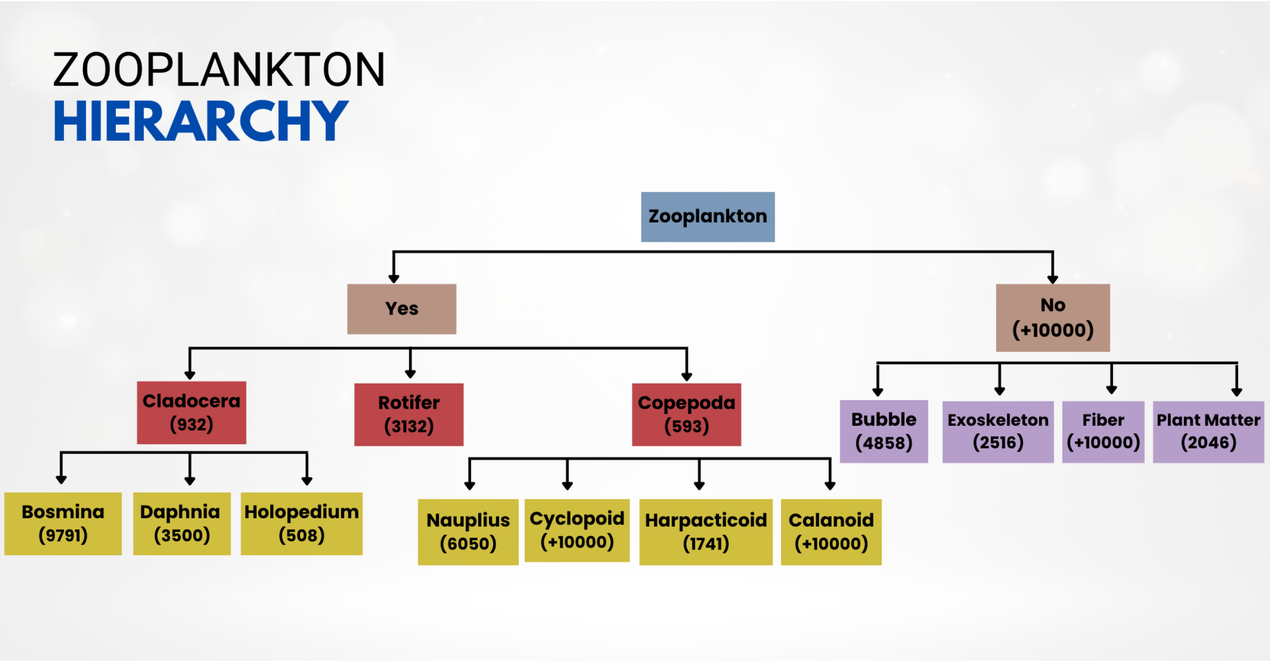

Model Architecture

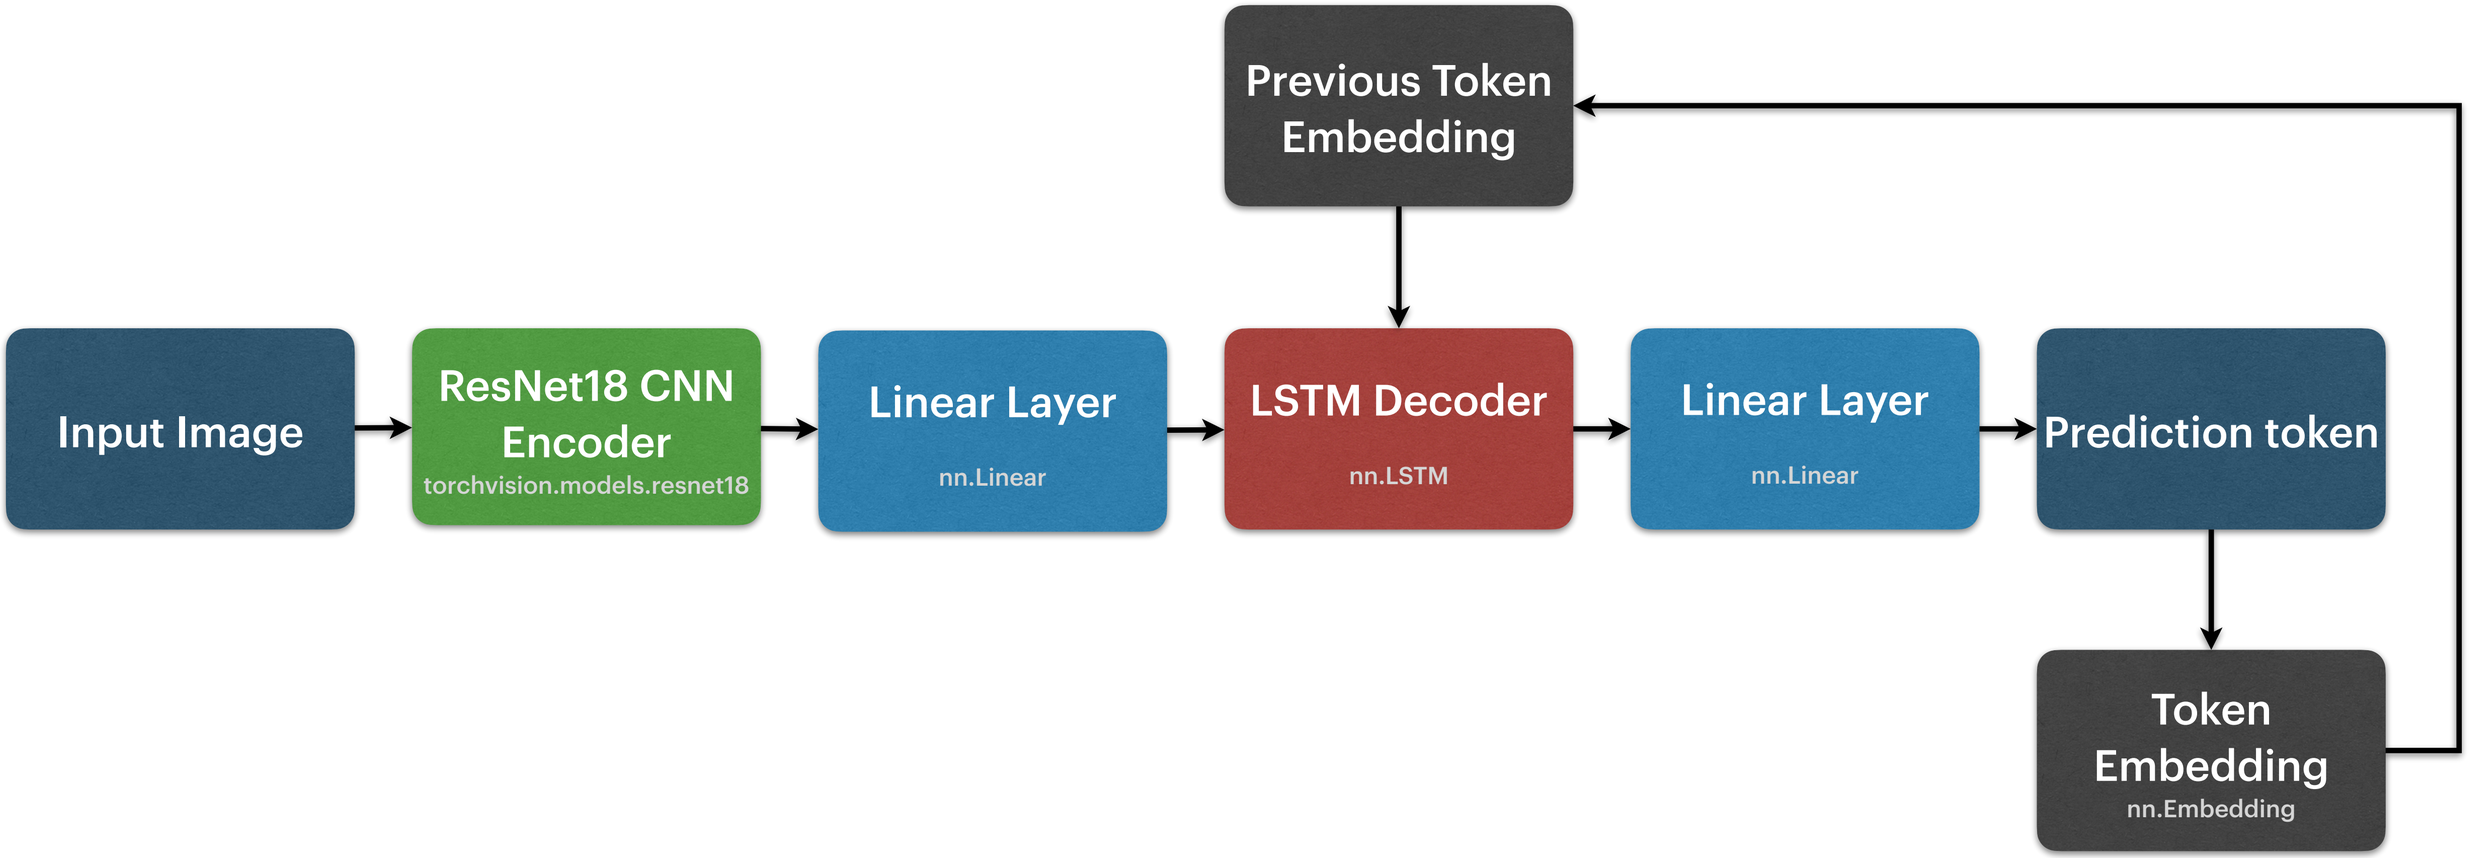

Example output: <START> → Root → Yes → Copepoda → Calanoid → <END>

Model Architecture

Why use this architecture ?

- **Follows naturally** from the problem

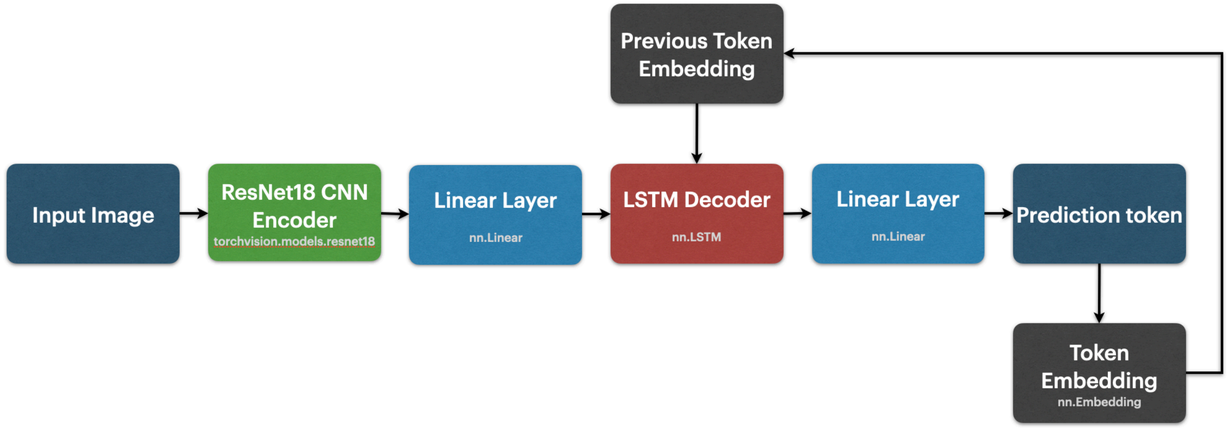

- Hierarchy can be **easily changed**

- Reuses LSTM module **reducing parameter count**

Training Loop

- 60/20/20 **Training/Val/Test**
- **Adam Optimizer** with **Gradient Clipping** max norm 1.0
- **Cross Entropy Loss** computed on flattened sequences jointly
- **Teacher Forcing Mode** feeds ground truth predictions as decoder outputs, making the model predict the next token
- **Dataset Augmentation** by rotating and flipping images for smaller classes

Parameter	Value
Embedding Size	512
Hidden Dimension Size	512
Learning Rate	0.0001
Max Epochs	50
Early Stopping Threshold	0.003
Early Stopping Patience	10 epochs

Tuned Hyperparameters

Classification Results

Leaf Level

- **95% Total Accuracy & F1 Score**
- **Best Predictions** on Daphnia (Cladocera), Nauplius (Copepod), Bubbles
- **Worst Predictions** on Harpacticoid (Copepod)

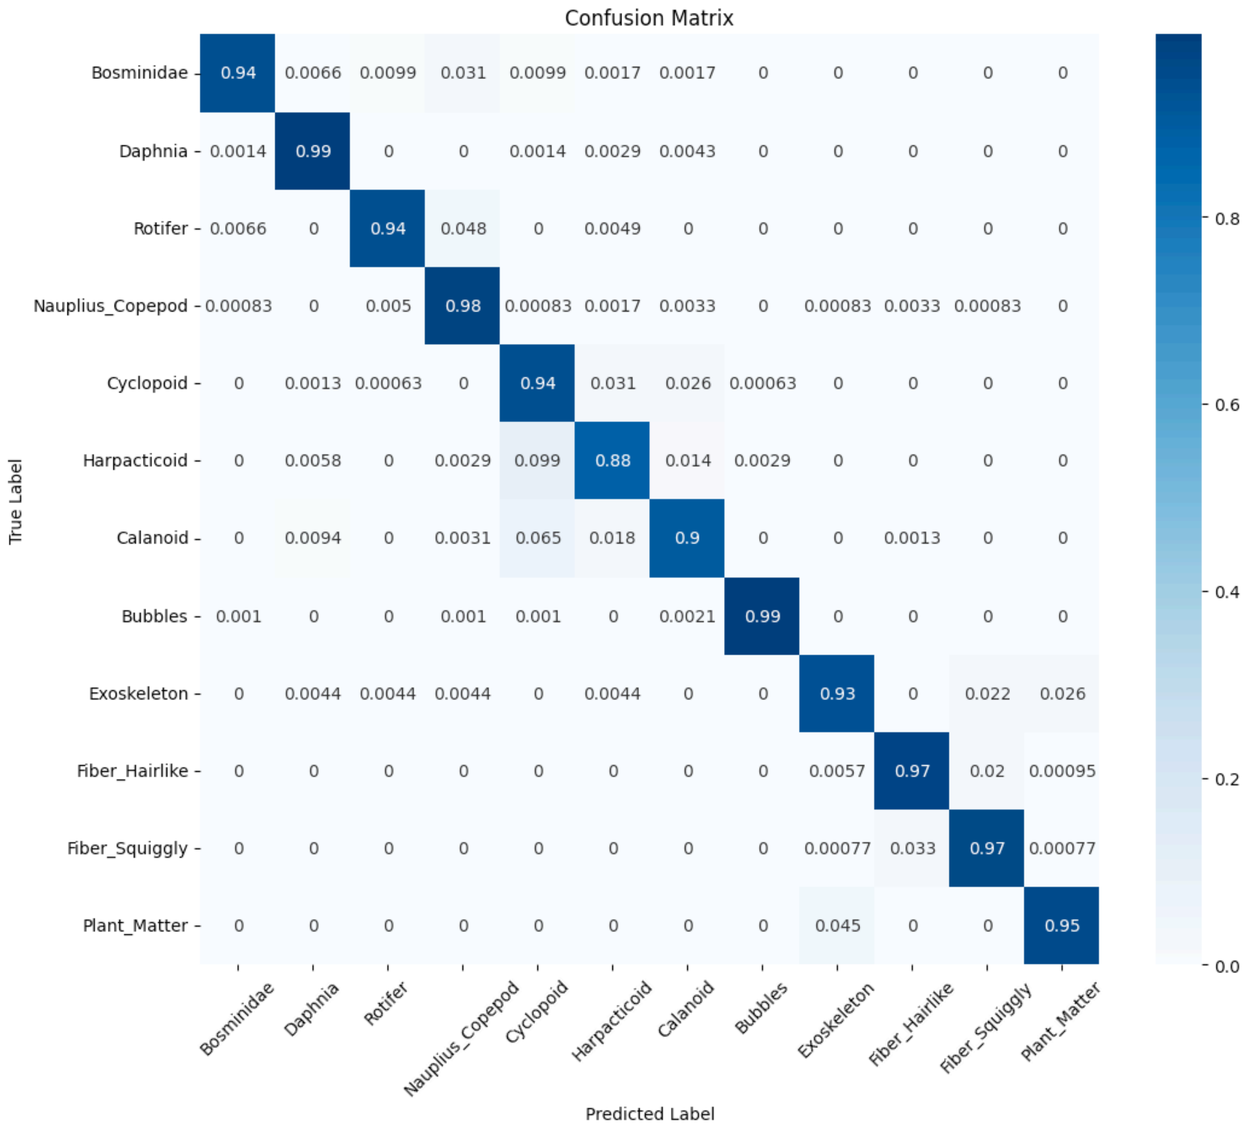

Classification Results

Parent Level

- **99% Total Accuracy & F1 Score**
- **Best Predictions** on Fiber, Copepoda, Bubbles
- **Worst Predictions** on Rotifer, Exoskeleton

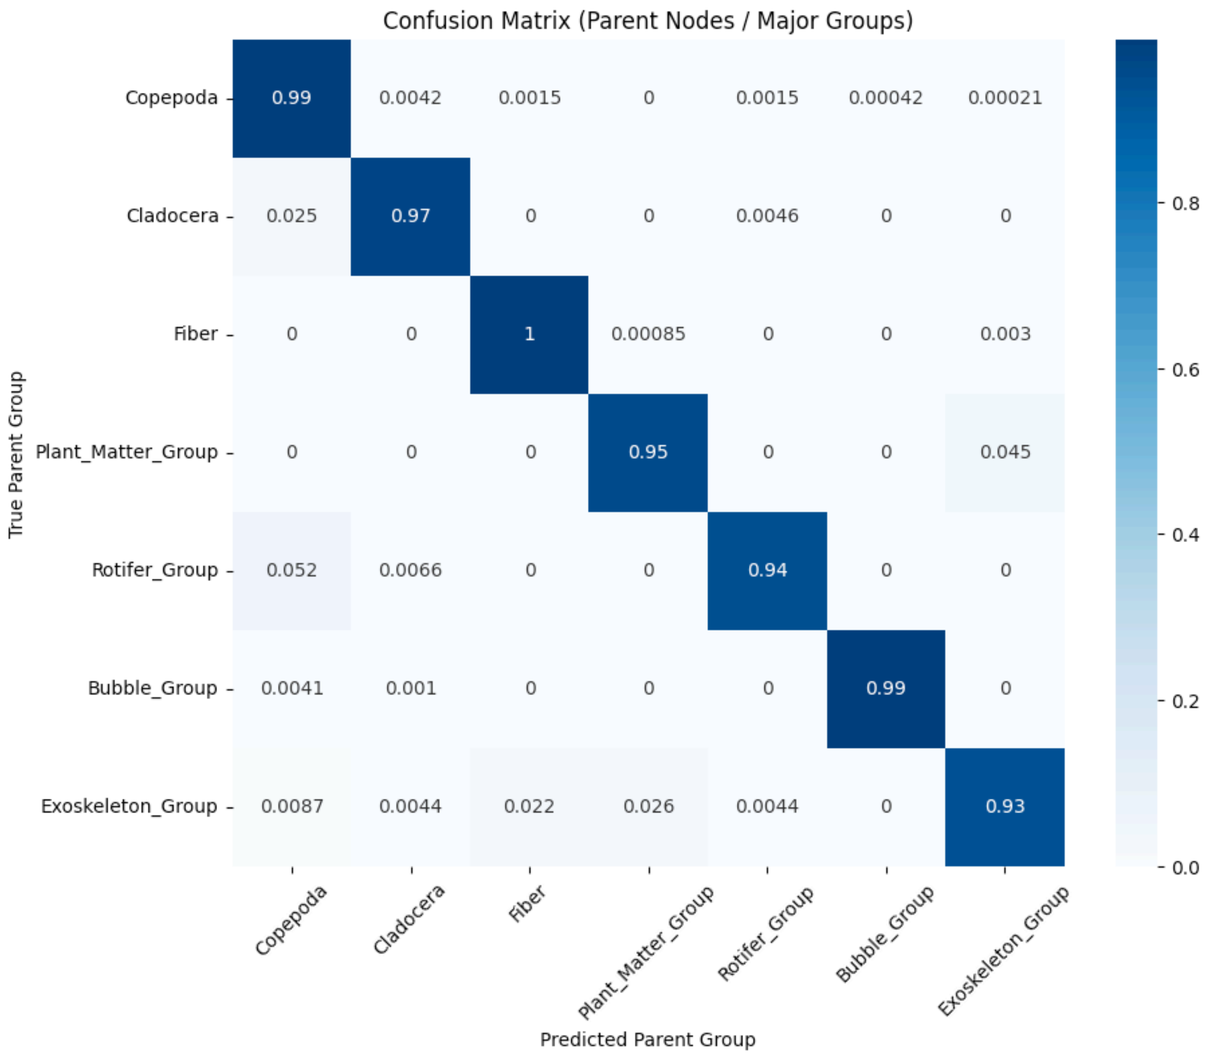

Thank You

Appendix

Frozen ResNet Results

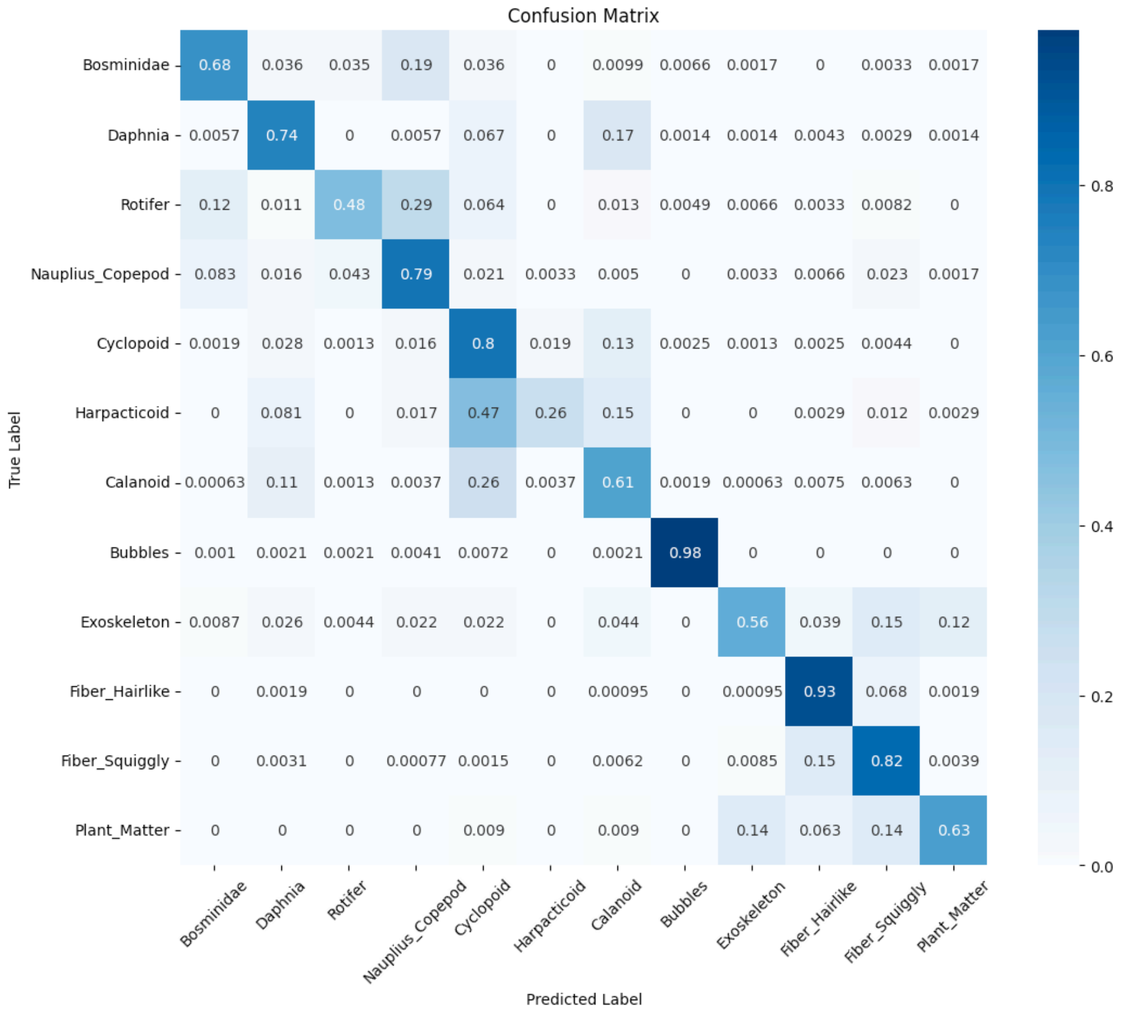

Leaf Level: 75% Accuracy

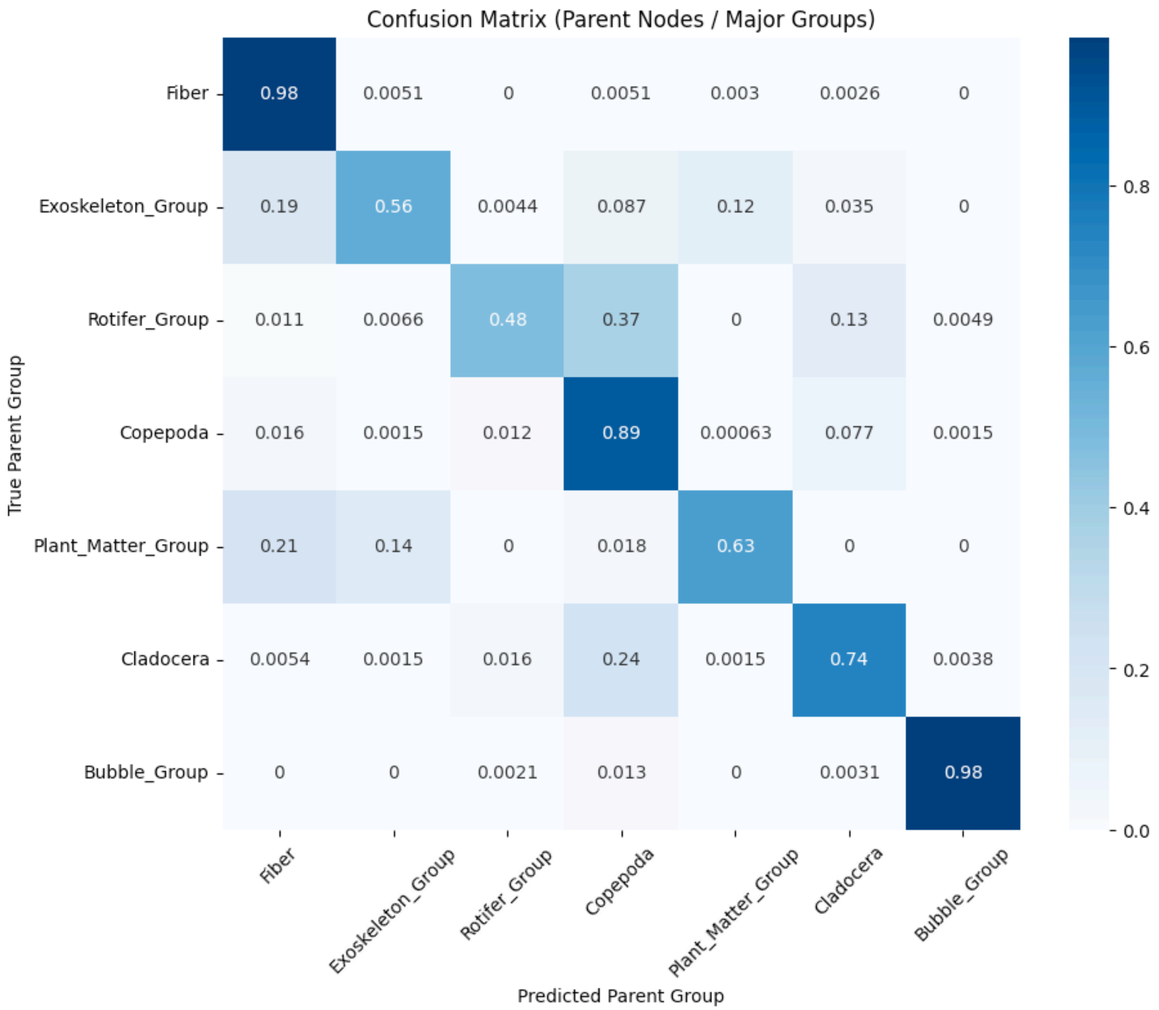

Parent Level: 87% Accuracy